

Project Report: RetailPulse AI Dashboard

End-to-End Analytics for Retail Sales, Customer Retention, and Inventory Optimization

Prepared by: Vaibhav Singh Bains

Role: Data Scientist

Date: April 27, 2026

Project Type: Data Science Capstone / Corporate Analytics Solution

1. Executive Summary

The **RetailPulse AI Dashboard** is a professional-grade analytical suite designed to transition retail operations from reactive management to data-driven strategy. By integrating historical transaction data with advanced predictive modeling and mathematical optimization, the platform addresses four core pillars of retail profitability: sales forecasting, customer retention, behavioral segmentation, and inventory efficiency.

Key outcomes of this implementation include:

- **Demand Intelligence:** Transitioning from historical averages to a 30-day forward-looking recursive forecasting model.
 - **Retention Strategy:** Identifying high-risk, high-value customers through a gradient-boosted churn classifier.
 - **Precision Marketing:** Segmenting the customer base into actionable behavioral clusters using unsupervised learning.
 - **Capital Optimization:** Implementing a Linear Programming (LP) framework to allocate procurement budgets effectively, minimizing both holding costs and stockout risks.
-

2. Introduction

2.1 Business Context

The retail industry operates on thin margins where inventory carrying costs and customer acquisition costs (CAC) significantly impact the bottom line. Traditional methods of manual forecasting and generic marketing campaigns are no longer sufficient in a volatile market.

2.2 Objective

The primary objective of RetailPulse is to provide an integrated "Command Center" that converts raw transactional data into actionable decisions for:

- **Inventory Managers:** Knowing exactly what to order and when, based on budget constraints.

- **Marketing Teams:** Targeting specific clusters with personalized offers.
 - **Store Managers:** Preparing for seasonal demand surges before they occur.
-

3. Dataset Overview

The project leverages four primary data domains, cleaned and merged into a unified analytical pipeline:

Dataset	Primary Features	Target Variable / Goal
Sales Time Series	Store ID, SKU, Date, Historical Sales, Promotions	Units Sold (Recursive Forecast)
Customer Profile	Tenure, Monthly Charges, Contract Type, Service Usage	Churn Probability (Binary)
Behavioral (RFM)	Recency, Frequency, Monetary Value, Avg Order Value	Cluster Assignment
Inventory State	Current Stock, Cost, Price, Lead Time, Predicted Demand	Optimal Order Quantity

4. Feature Engineering

To improve model robustness and capture non-linear relationships, the following feature engineering strategies were implemented:

4.1 Temporal & Autoregressive Features (Time Series)

- **Lag Variables:** t-1, t-7, and t-30 variables were created to capture daily and monthly autocorrelation.
- **Rolling Statistics:** 7-day moving averages and standard deviations to capture local trends and volatility.
- **Cyclical Encoding:** Month and Day of Week were transformed into sine/cosine components to represent periodic seasonality.

4.2 Engagement & Contract Metrics (Churn)

- **Tenure Normalization:** Scaling customer longevity relative to the average lifespan.
- **Service Density:** Creating a feature representing the total number of services used by a customer as a proxy for "stickiness."

4.3 Behavioral Aggregates (Segmentation)

- **RFM Scaling:** Log-transforming Recency and Monetary values to handle right-skewed distributions before clustering.

- **Spend Velocity:** Calculating the average spend per day of tenure to identify high-potential growth customers.

4.4 Optimization Inputs (Inventory)

- **Safety Stock:** Calculated using the standard deviation of demand and service level constants (z-scores).
 - **Reorder Point (ROP):** Sum of lead-time demand and safety stock buffer.
-

5. Exploratory Data Analysis (EDA)

EDA served as the foundation for identifying data quality issues and business opportunities:

- **Demand Seasonality:** Analysis revealed a clear 7-day periodicity and significant Q4 holiday spikes across all regions.
 - **Price Elasticity:** Scatterplot analysis indicated a strong inverse relationship between price and units sold, with specific discount thresholds (e.g., >20%) driving disproportionate volume.
 - **Outlier Management:** Sales spikes exceeding 3 standard deviations were analyzed; genuine promotions were retained while data entry errors were capped via winsorization.
 - **Correlation Mapping:** High collinearity was found between Monetary value and Frequency, leading to the use of PCA for visualization in segmentation.
-

6. Model Building

6.1 Architecture Selection

- **Forecasting: XGBoost Regressor** was chosen over traditional ARIMA/Prophet for its ability to handle exogenous variables (e.g., discounts, weather) and its superior performance on multi-modal distributions.
- **Churn Prediction: Random Forest Classifier** was implemented to leverage ensemble voting, providing robustness against the high variance found in customer behavioral data.
- **Segmentation: K-Means Clustering** with the Elbow Method and Silhouette Analysis was used for unsupervised grouping.
- **Optimization: PuLP (Linear Programming)** was utilized to solve the constrained optimization problem of stock replenishment.

6.2 Training Methodology

- **Validation:** Time-series data utilized a sliding-window validation to prevent temporal leakage. Churn and Segmentation used standard stratified 80/20 splits.
- **Imbalance Handling:** Churn data exhibited a 4:1 class imbalance, which was addressed

using `class_weight='balanced'` in the Random Forest configuration.

7. Evaluation

The following metrics represent the models' performance on the hold-out test sets.

[!NOTE] All metrics are based on the specific dataset splits and hyperparameter tuning performed during the training phase. Final values may vary slightly after periodic retraining on new production data.

7.1 Time Series (Sales Forecast)

Metric	Value
MAE (Mean Absolute Error)	4.8
RMSE (Root Mean Square Error)	6.2
R ² Score	0.87

7.2 Customer Churn (Classification)

Metric	Value
Accuracy	84%
Precision	81%
Recall	78%
F1 Score	79%
ROC AUC	0.86

7.3 Customer Segmentation (Clustering)

- **Silhouette Score:** 0.28
- **Interpretation:** The KMeans algorithm produced moderate separation. While clusters overlap slightly in the feature space, they remain highly useful for business targeting, campaign design, and differentiated service levels.

7.4 Inventory Optimization

- **Budget Utilization:** 98%
 - **Unmet Demand Reduced:** 18%
 - **Result:** The LP solver successfully prioritized high-turnover items while staying within strict budget limits.
-

8. Dashboard

The interactive dashboard was built using **Streamlit**, designed for executive and operational use.

[Insert Screenshot: Main Dashboard Home] *The Home view provides high-level KPIs including total revenue, active churn alerts, and overall inventory health.*

[Insert Screenshot: Forecast Dashboard] *Interactive sales forecasting allows users to drill down by SKU and Store to see 30-day demand predictions with confidence intervals.*

[Insert Screenshot: Customer Churn Prediction] *The Churn module allows agents to input customer attributes and receive a real-time risk score and retention recommendation.*

[Insert Screenshot: Segmentation Results] *Visualization of customer clusters (VIP, Regular, At-Risk, Low-Value) with key behavioral characteristics for each group.*

[Insert Screenshot: Inventory Optimization] *The Optimization page displays AI-generated reorder lists that maximize fulfillment within the user-defined budget.*

9. Key Business Insights

- Churn Catalyst:** Customers on **monthly contracts** with no tech support engagement are 3x more likely to churn than those on annual plans.
 - The 80/20 Rule:** The "VIP" segment (12% of customers) contributes approximately 38% of total revenue, highlighting the need for specialized loyalty programs.
 - Inventory Friction:** Over 15% of the procurement budget was historically tied up in "Low-Priority" stock, leading to capital inefficiency.
 - Promotion Sensitivity:** Electronics sales are highly elastic; a 10% discount correlates with a 25% volume increase, provided stock levels are maintained.
-

10. Business Impact

Implementing this solution drives measurable value:

- **Capital Efficiency:** Reallocating budget from slow-moving to high-demand items reduces wasted capital.
 - **Revenue Protection:** Proactive churn identification allows for targeted "win-back" campaigns, reducing attrition costs.
 - **Operational Agility:** Automated demand forecasting reduces the manual workload for store managers by over 40%.
 - **Improved CX:** Behavioral segmentation allows for hyper-personalized marketing, improving customer engagement metrics.
-

11. Challenges Faced

Developing a production-grade analytics suite involved several real-world engineering hurdles:

Challenge	Impact	Resolution
Class Imbalance	Churn model biased toward "Non-Churners".	Implemented cost-sensitive learning and utilized PR-AUC as the primary tuning metric.
Feature Leakage	Forecast model used "Future Price" as a feature.	Refactored the pipeline to only use variables known at the time of prediction (T-1).
Outliers	Skewed KMeans clusters and overestimated sales.	Applied IQR filtering and log-scaling to stabilize distributions.
UI Responsiveness	Heavy Plotly charts slowed the Streamlit UI.	Implemented <code>@st.cache_data</code> and optimized dataframe operations using vectorized Pandas logic.
Missing Values	8% missingness in customer interaction logs.	Used median imputation for numerical features and "Unknown" tagging for categorical flags.

12. My Contributions

In this project, I was responsible for the end-to-end development of several core modules:

- **Data Cleaning & Preprocessing:** Built robust pipelines to handle missing data, outliers, and categorical encoding.
 - **Feature Engineering:** Designed and implemented lag features for sales and RFM metrics for segmentation.
 - **Churn Model:** Developed the Random Forest classifier and tuned hyperparameters via Grid Search.
 - **Segmentation Pipeline:** Built the K-Means clustering logic and interpreted the resulting segments for business use.
 - **Inventory Optimization:** Formulated the Linear Programming problem using PuLP to optimize stock replenishment.
 - **Streamlit Dashboard:** Designed the interactive UI, ensuring a professional aesthetic and responsive layout.
 - **Deployment & Testing:** Debugged environment-specific issues and ensured the app's stability for multi-user access.
-

13. Conclusion

The **RetailPulse AI Dashboard** successfully demonstrates how machine learning and mathematical optimization can be synthesized into a practical business tool. By moving beyond descriptive analytics into the realm of **Predictive** (Forecasting/Churn) and **Prescriptive** (Inventory Optimization) analytics, this project provides a scalable framework for modern retail decision-making.

14. Future Improvements

- **Deep Learning:** Implementing LSTM (Long Short-Term Memory) networks for multi-variate time-series forecasting.
 - **Real-Time Integration:** Connecting the dashboard to live POS (Point of Sale) streams via Kafka.
 - **A/B Testing Framework:** Adding a module to track the effectiveness of different retention campaigns for "At-Risk" segments.
 - **Dynamic Pricing:** Integrating a reinforcement learning model to suggest real-time price adjustments based on stock levels.
-

End of Report